

Using Hadoop and Spark for Big Data Analytics: Predicting Readmission of Diabetic patients

KUMAR SATYAM^{1,*}, PIYUSH SHINDE^{1,**}, AND SRIKANTH RAMANAM^{1,***}

¹School of Informatics and Computing, Bloomington, IN 47408, U.S.A.

* Corresponding authors: ksatyam@indiana.edu

** Corresponding authors: pshinde@iu.edu

*** Corresponding authors: srikrama@iu.edu

project-000, April 28, 2017

This project proposes and demonstrates the use of Hadoop and Spark on cloud to run predictive analytics using machine learning on large amount of data. Our case study is to predict the readmission likelihood for diabetes patients using their available medical history.

© 2017 <https://creativecommons.org/licenses/>. The authors verify that the text is not plagiarized.

Keywords: Hadoop, Spark, MLlib, Ansible, Cloudmesh Client, Predictive Analysis

<https://github.com/cloudmesh/classes/blob/master/project/S17-IR-P004/report/report.pdf>

CONTENTS

1	Introduction	1
2	Architecture	2
3	Technologies	2
4	Cloud Infrastructure	3
4.1	Chameleon Cloud	3
4.2	Jetstream Cloud	3
4.3	Virtual Box	3
5	Automated Cloud Deployment: Ansible	3
6	Data Cleansing and Pre-processing	4
7	Preliminary Data Analysis	5
8	Data Analysis on Spark cluster	5
8.1	Start the Spark service	5
8.2	Data Storage	5
8.3	Launching Data Analysis Application	6
9	Results	6
10	Benchmarking	6
10.1	Computation time of algorithms in clouds	6
10.2	Comparison of multiple algorithms in Spark MLlib vs scikit-learn	6
10.3	Comparison of Spark Deployment Time	7
10.4	Comparison of File Transfer Time	7

11	Troubleshooting	7
12	Conclusion	8
13	Acknowledgments	8

1. INTRODUCTION

The idea behind this project is to introduce Hadoop/Spark over cloud infrastructure as a scalable and faster solution for predictive analysis using machine learning.

We chose the case study of predicting the likelihood of a diabetes patient getting readmitted within 30 days from the date of discharge using his/her available medical data. We approached this problem as a classification problem to classify the patients into 'Yes' or 'No' classes, indicating whether the patient is likely to be readmitted or not in the next 30 days. We used different machine learning algorithms on the available data, after some pre-processing, to predict the same. The accuracy percentages obtained for all the utilized classification algorithms are included in the report. We also compared the results of Spark's Mlib algorithms against scikit-learn's algorithms and found that both yielded similar results.

We approached the solution from a pure data perspective to address the challenge of lacking medical domain knowledge. For some basic information, we relied on the dataset description on UCI website [1], ICD-9 [2] and earlier studies [3]. We followed a workflow as shown in figure 1.

Our other important goal is to propose an end-to-end solution that is scalable and faster. While our dataset is about 100,000 records, anticipating real-world scenarios with huge amounts of data we are proposing a Hadoop based solution. We are

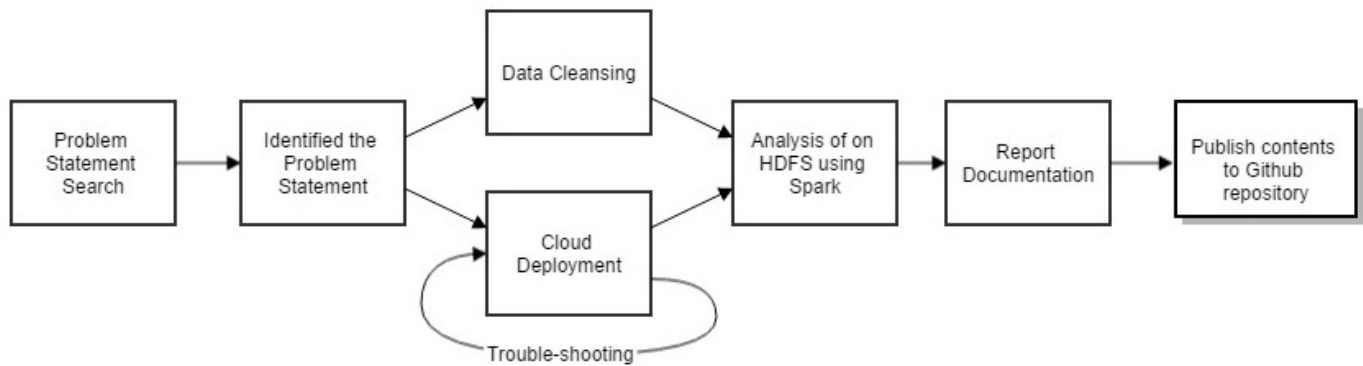


Fig. 1. Workflow of the project

utilizing Spark for its faster processing [4] and advanced analytics through packages like Mllib [5], which provides several commonly used machine learning algorithms. Finally we are implementing this solution over cloud infrastructure to meet our infrastructure requirements. This helped us demonstrate an end-to-end solution closer to real-world scenarios, where enterprises utilize cloud infrastructure in a pay per-use model. This helped us save time and resources in setting up the infrastructure.

We deployed our solution on two different clouds and obtained metrics to assess the infrastructure performance with our solution. We also deployed our solutions on a distributed Hadoop/Spark environment built on on-premise machines. We compared the performance metrics of all the three infrastructure choices, with respect to our solution and included them in this report.

2. ARCHITECTURE

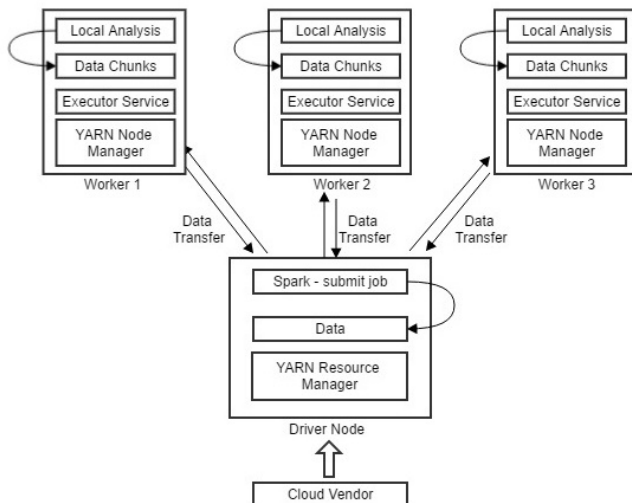


Fig. 2. Architecture Diagram

Figure 2 gives an overview of our solution's architecture. We deployed a spark driver node and three worker nodes. A driver node is a node that runs the driver program. It declares the transformations and actions on RDDs (Resilient Distributed Datasets) of data and submits such requests to the master [6]. In practical terms, the driver is the program that creates the Spark

Context [7], connecting to a given Spark Master. It is a node where the yarn resource manager resides. A worker node is a node, which executes the program that involves individual data analysis task. Running the spark-submit script from the master node starts the spark job. It divides the data into data chunks and transfers them to individual worker nodes. Then a processing task is performed on the data chunks on the individual worker nodes. The processed data and analytics results are then written back to the HDFS file system as needed.

We then deployed Hadoop and Spark on these machines to setup a HDFS Spark cluster on our cloud and on-premise machines. We stored our dataset on the Hadoop's HDFS file system. Finally, we ran our predictive analytics application, which utilizes Mllib, on Spark cluster by launching it using spark-submit.

3. TECHNOLOGIES

Technology	Usage
Hadoop [8]/ Spark [9]	Big Data Technologies
Python [10]	Development
Mllib [5]/ scikit-learn [11]	Machine Learning Library
GitHub [12]	Project Repository
Ansible [13]	Application Deployment & Configuration Management
Chameleon [14], JetStream [15], VirtualBox [16]	Benchmarking
LaTeX [17]	Document Preparation

Table 1. List of technologies used

We used specific technologies for specific tasks in this project as listed in table 1:

- **Hadoop:** Apache Hadoop is a framework for processing and storing huge amounts of data, commonly known as 'Big Data', in a distributed applications. It allows users to build scalable and highly available data applications. Hadoop has four modules: HDFS, YARN, MapReduce and Hadoop Common. Hadoop Distributed File System (HDFS) allows

users to store large amounts of data. YARN is the framework for job scheduling and resource management. MapReduce supports parallel processing of large data sets stored in the distributed environment through HDFS. Hadoop Common provides utilities that support other Hadoop modules.

- **Spark:** Spark runs on Hadoop and provides faster data processing capabilities for data on HDFS. It primarily uses a new data structure called Resilient Distributed Dataset (RDD) for processing. RDD is a read only multiset of data items distributed over cluster of virtual machines. Spark also has a new feature of fault tolerance and in the event of a primary master node failure, the secondary master takes over. Spark, unlike Hadoop applications, allows the iterative reading and writing in-memory. After processing it writes the data to HDFS.
- **Python:** We chose python as per our programming language. Python is one of the programming languages supported by Spark API through pyspark. We used it because of its simple syntax and data manipulation capabilities.
- **scikit-learn:** It is an open source Python library that provides Machine Learning algorithms and other utilities to preprocess and visualize data [11].
- **MLlib:** Spark MLlib is the Spark's machine learning library provides machine learning algorithms that can be applied on Resilient Distributed Datasets [18]. It also provides other data manipulation utilities. MLlib has API available in Java, Python, Scala and R [5].
- **GitHub:** GitHub is 'a web-based Git or version control repository and Internet hosting service' [19]. We used Github repositories to store all the files related to documentation, ansible scripts and python code.
- **Ansible:** We are using ansible for automating deployment of our software on cloud. We used ansible scripts to automate deployment of cloudmesh client along with its pre-requisites like pip, virtualenv etc. We also used ansible for automating deployment of Hadoop and Spark.

4. CLOUD INFRASTRUCTURE

We have setup the required infrastructure by provisioning virtual machines on two cloud vendors, Chameleon, Jetstream and our on-premise machines.

4.1. Chameleon Cloud

Chameleon is a collaborative cloud service primarily meant for research community. It allows users to explore problems ranging from the creation of Software as a Service to kernel support for virtualization. It is a good example of IAAS loaded with software defined networking and optimized virtualization technologies. We created three virtual machines on this cloud. One for Master node and two for worker nodes of Spark.

4.2. Jetstream Cloud

Jetstream is a cloud service which aims to provide researchers Jetstream's development was led by Indiana University's Pervasive Technology Institute (PTI) in collaboration with other universities [15] across the United States. This cloud service was used to provision the necessary virtual machines. We created three virtual machines on Jetstream for Spark cluster nodes.

4.3. Virtual Box

It is a fully virtualized hypervisor which gives the ability to spawn virtual machines in local commodity hardware. In fully virtualized environment, the guest OS is not aware of the underlying resources on which it is running as the hypervisor creates a complete simulation of the underlying hardware.

A brief comparison of multiple attributes of the clouds used are displayed in table 2.

Clouds	Chameleon	Jetstream	VirtualBox
CPU	Intel Xeon X5550	Dual Intel E-2680v3 "Haswell"	Intel Core i5-6200U
RAM	4 GB	2 GB	2 GB
Number of CPU's	1	1	1
CPU Cores	1	1	2
CPU Speed	2.3 GHz	2.3 GHz	2.3 GHz

Table 2. Comparison of cloud vendors

5. AUTOMATED CLOUD DEPLOYMENT: ANSIBLE

For this project, we used Ansible to automate the deployment of spark and its prerequisites. The Ansible script is written such that we can leverage the cloudmesh client technology to deploy the spark cluster. The steps involved in the script can be seen in figure 3.

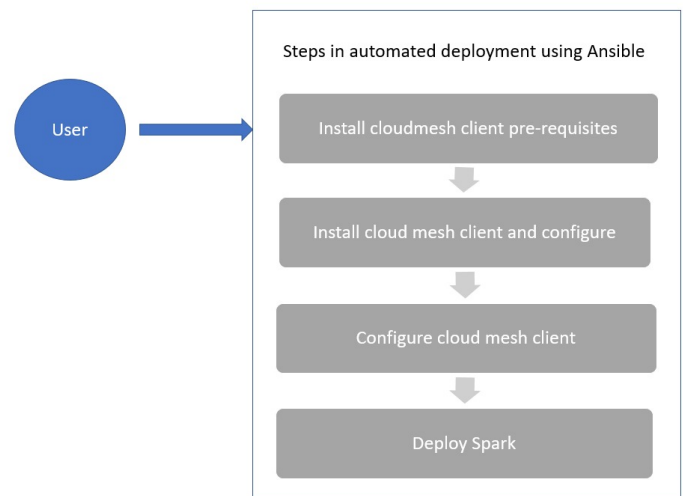


Fig. 3. Automated Cloud Deployment using Ansible

The Ansible playbook package constitutes 4 files.

1. *inventory*: This file is used by the playbook to install packages and services on target servers. This Inventory file consists of target servers and these servers can be categorized into groups like web servers, db servers etc. Under each group, we can list the server which serves common purpose.
2. *playbook-cloudmesh-first-time-install.yml*: This file was used for deployment of cloudmesh client.

3. *host*: This file contains the list of hosts.
4. *hadoop-spark-playbook.yml*: This consists of the ansible code for redeployment of all the pre-requisites packages for clouddmesh client. This also deployment of spark using clouddmesh client.

The following procedure has to be followed to run the Ansible script:

1. Install clouddmesh client for the first run 'playbookclouddmesh-first-time-install.yml' with the following command. This automates the deployment of clouddmesh client, which is a pre-requisite for installing Spark over our cloud infrastructure.

```
ansible-playbook
playbook-clouddmesh-first-time-install.yml
--ask-sudo-pass -vvvv
```

2. Open, /.clouddmesh/clouddmesh.yml and edit the following section, the entry with <> should be customized as per your credentials.

```
profile:
  firstname: <first name>
  lastname: <last name>
  email: <email id>
  user: <chameleon/jetsream/other cloud username>
```

3. Change the entry of active cloud to use your preferred cloud. Example with chameleon is present below:

```
active:
  - chameleon
clouds:
  ...
```

4. Under a cloud(chameleon/jetstream/..) change the following entry, the entry with <> should be customized as per your credentials:

```
credentials:
OS_PASSWORD: <enter your chameleon cloud password here>
OS_TENANT_NAME: CH-818664
OS_TENANT_ID: CH-818664
OS_PROJECT_NAME: CH-818664
OS_USERNAME: <username>
```

5. Also change, the type of OS and flavor you want for hadoop/spark cluster provided by your preferred cloud:

```
default:
  flavor: m1.medium
  image: CC-Ubuntu14.04
```

6. Edit the file hadoop-spark-playbook.yml in the section: 'name:Preparing clouddmesh- setting default user', the entry with <> should be customized as per user credentials:

```
- name: Preparing clouddmesh- setting default user
  become_user: "{{ lookup('env','USER') }}"
  shell: cm default user=<chameleon/jetstream user>
```

7. Now run the following command to deploy hadoop/spark cluster which will run the hadoop-spark-playbook.yml. It will again upgrade clouddmesh client, so it is up to date:

```
ansible-playbook hadoop-spark-playbook.yml
--ask-sudo-pass -vvvv
```

6. DATA CLEANSING AND PRE-PROCESSING

The initial data set is publicly available on the UCI Machine Learning Repository [1]. The initial data set has information extracted from the database satisfying the following criteria [3]:

1. Each row corresponds to an inpatient encounter (a hospital admission).
2. All of the encounters are "diabetic" encounters, that is, one during which any kind of diabetes was entered to the system as a diagnosis.
3. Each encounter also corresponds to a patient stay between 1 and 14 days.
4. Laboratory tests were performed during the encounter.
5. Medications were administered during the encounter.

101,766 encounters were present in the data set that satisfy the above five inclusion criteria. Each encounter consists of 55 features describing the diabetic encounters, including demographics, diagnoses, diabetic medications, number of visits in the year preceding the encounter, and payer information. We defined the readmission field with 2 values: "YES," for cases where the patient was readmitted within 30 days of discharge and "NO," for both readmission after 30 days scenario and no readmission at all.

Diagnosis 1, 2 and 3 had many categorical values in the form of ICD-9 codes and had missing values. These ICD-9 codes were sorted and grouped into 9 categories, namely Circulatory, Respiratory, Digestive, Diabetes, Injury, Musculoskeletal, Genitourinary, Neoplasms and Other based on the ICD-9 codes [3]. The missing values were assigned the group 'Other'.

This data set had several features with empty fields. So we removed the features missing high percentages of data as they affect our analysis. The removed features were weight, medical specialty and payer code. The race attribute had 2% missing values which were filled by the mode value 'Caucasian'.

Observations only with unique patient ids were considered, excluding those with discharge disposition corresponding to the patient's death.

We filtered our data set according to the above-mentioned constraints and retained 62,937 encounters each corresponding to an unique patient and 55 features describing such encounter. We prepared three data-sets to test the multiple machine learning algorithms (Stochastic Gradient Descent, Gaussian Naïve Bayes, K-means and Decision Tree). Finally we also removed the patient id and encounter id as they were not relevant to learning algorithms.

We used one hot encoding to convert the categorical data features to numerical data. This creates new dummy features to represent the categorical data in numerical format.

The first data set was prepared using one hot encoding on the original data (with 62,937 observations), that resulted in 136 features representing the original 55 features.

For our second data set we implemented feature selection using a Variance Threshold algorithm that removes all low-variance features [20]. We set a variance threshold of '0.8'.

The data set B was formed using this algorithm, which helped extract 26 features.

The original data contains the age attribute grouped in 10-year intervals from 0 to 100 years. For the third data set we grouped the 10 intervals to 3 intervals by combining the age groups younger than 30, 30-60 and older than 60 years. One-hot encoding was then applied to form the third data set. This data set contained 129 features.

7. PRELIMINARY DATA ANALYSIS

The unit of our analysis is an encounter; to keep the observations independent, we only analyzed one encounter per patient. We performed early data analysis in python in a local machine. We implemented four classification algorithms using scikit-learn library on each of the 3 data sets created. The data sets were first divided in training and testing set. The training set had about 80% observations (5000 observations approx.) whereas the testing set had the remaining 20% observations (12937 observations). Each of the algorithms provided by scikit-learn were used following in the manner:

1. Create and fit a model using the observations and readmission of the training set.
2. Predict labels of the testing set.
3. Calculate the accuracy using the predicted labels and true labels of the testing set. The parameters needed in implementing the above-mentioned algorithms were set so as to be valid to our data, give optimum results and make the results to be reproducible.

KMeans clustering gave us approximately 55% accuracy with all the three data sets. Though it is an unsupervised learning algorithm, we used it to examine if the clustering divided the data into readmission classes, Yes and No, to an acceptable level. We concluded that clustering based on the Euclidean distance may not be the right approach for this classifying this data set. The accuracy percentages obtained for other classification algorithms can be found in the Table 3.

8. DATA ANALYSIS ON SPARK CLUSTER

We performed data analysis on Spark cluster using pyspark MLlib on data stored on HDFS.

8.1. Start the Spark service

Start the service of spark using the following command:

```
$SPARK_HOME/sbin/start-all.sh
```

Stop the service of spark using the following command:

```
$SPARK_HOME/sbin/stop-all.sh
```

Using the command 'jps' we get a list of the following services:

```
nodemanager
resourcemanager
master
namenode
applicationmaster
```

Classification Technique	Number of Features	Accuracy (%)
SGDClassifier	136	90.68
	26	86.31
	129	86.96
GaussianNB	136	10.43
	26	88.30
	129	9.96
KMeans	136	55.26
	26	55.24
	129	55.26
DecisionTreeClassifier	136	83.35
	26	82.50
	129	83.28

Table 3. Results from scikit-learn

8.2. Data Storage

Uploading input data and code to Driver Node After the Spark setup is ready for the deployment, the data is pushed from the localhost to the remote spark master node. For this we used an ansible script.

1. Ansible script

- (a) In the host files we set the target master(driver node) IP address as follows:

```
[remotehosts]
129.114.33.106 ansible_ssh_user=cc
```

- (b) Now , add the following entry in the yaml file to transfer the file to destination

```
- hosts: remotehosts
  tasks:
    - name: Transfer file from local to
      satyam-001
      synchronize:
        src: /home/<username>/ansible
        script/ansible-spark/trainindat.csv
        dest: /home/cc/
        mode: push
        delegate_to: 127.0.0.1
```

2. Using Github we uploaded the input data csv file and python execution code to a git repository. We installed git package in the spark driver node. We used 'wget' command with the repository path to download the data set to the virtual machine.

After the data is downloaded to the virtual cloud machine, we uploaded the file to HDFS through the following command.

```
Hdfs dfs -put <source file path> <hdfs-folderpath>
```

This HDFS file serves as an input for our analytics application.

8.3. Launching Data Analysis Application

We used python programming language to develop an application that performs predictive analytics tasks. We leveraged pyspark.mllib library for machine learning algorithms.

We launched our application code using the following command

```
$ ./bin/spark-submit --class path.to.your.Class
--master yarn --deploy-mode cluster [options]
<app jar> [app options]
```

There are 4 main steps to each implementation of.

- Input formatting: MLib classes expect RDD's of Labeled-Points. For this we parsed the data and converted each entry into a LabeledPoint, with label specifying the true output class.
- Next, the processed data frame is divided into train and test datasets.
- Train the model with the algorithm and training data
- After training, We used the model to predict the classes for test data and calculate the accuracy

9. RESULTS

Classification Technique	Number of Features	Accuracy (%)
SGDClassifier	136	90.88
	26	90.99
	129	90.88
GaussianNB	136	84.97
	26	85.02
	129	84.987
KMeans	136	55
	26	55
	129	55
DecisionTreeClassifier	136	91.01
	26	90.75
	129	90

Table 4. Results from Spark MLib

Kmeans did not successfully separate the data class wise as observed in scikit-learn library, as shown in table 4. We used supervised classification algorithms namely decision tree, logistic stochastic gradient descent and naïve bayes. We obtained good accuracies of 80-90% with classification algorithms.

These can be found in the below table. Similar analysis in a real world scenario can be utilized to predict the readmission likelihood using medical records of diabetes patients. This enables doctors to pay special attention to those patients, identify the causal factors and give preventive care.

10. BENCHMARKING

In the figure 4 we are displaying the time taken to run each of the four algorithms in different clouds. The four algorithms are SGD Classifier, Gaussian Naïve Bayes, kmeans and Decision Tree Classifier.

1. The complexity of the algorithm. For example, kmeans clustering has $O(n \log n)$ time complexity which is worse than Gaussian Naïve Bayes and hence takes more time.
2. The compute resources used. For example the chameleon cloud takes less time to execute any other algorithm because the RAM configuration for the VMs of chameleon cloud is better than any other VMs.
3. The network latency between the hosts: The VMs which are provisioned on clouds can be on different hosts spread across different racks of datacenters and even datacenters across geographies. This may result in network latency and affect the runtime of a program running in a distributed environment.

10.1. Computation time of algorithms in clouds

Figure 4 shows the performance of multiple machine learning algorithms using Spark MLib on clusters in Jetstream, Chameleon and Virtual Box.




Fig. 4. Run-time in different clouds.

Figures 5 and 6 shows the run-time comparison for machine learning algorithms using Spark MLib in Jetstream and Chameleon, for predicting readmission likelihood, for 3 and 5 VM's respectively.




Fig. 5. Run-time in different clouds for 3 vms

10.2. Comparison of multiple algorithms in Spark MLib vs scikit-learn

We compared the accuracies for different MLib algorithms using pyspark.mllib package. We can see from figure 7 the accuracies




Fig. 6. Run-time in different clouds for 5 vms

achieved for multiple algorithms in MLlib. While Kmeans gave the lowest accuracy, rest of the algorithms gave similar accuracies.




Fig. 7. Accuracy Results with pyspark.mllib

We compared the accuracies achieved from Spark MLlib and scikit-learn algorithms. Figure 8 shows accuracy results of multiple machine learning algorithms in scikit-learn vs MLlib.




Fig. 8. Accuracy Results: scikit-learn vs pyspark.mllib

10.3. Comparison of Spark Deployment Time

We compared the Hadoop-Spark deployment time with different number of machines on each of the clouds, Jetstream and Chameleon. We can see from figure 9 that the Hadoop-Spark deployment is faster on Chameleon than Jetstream.

10.4. Comparison of File Transfer Time

We uploaded two files of different sizes to both Jetstream and Chameleon. We recorded the file transfer times for each of these




Fig. 9. Comparison of Spark deployment results

clouds. We can see from figure 10 that it took lesser time to transfer files JetStream cloud machines than to Chameleon cloud machines.




Fig. 10. Comparison of file transfer time

11. TROUBLESHOOTING

We encountered several errors while running different commands in the Linux terminal. We have listed some of the commands and errors encountered while executing them followed by the steps to resolve them.

1. Command: cm ssh key upload

Error 1: Permission denied (public key)

Error 2: Problem uploading key <username> to cloud chameleon: The request you have made requires authentication

Resolution:

```
cm ssh-key delete <username>-001
cm ssh-key delete <username>-001 cloud = chameleon
ssh-add
cm key add --ssh
cm ssh-key upload
cm key refresh
```

2. Command: cm hadoop deploy

Error: IndexError: list index out of range

Resolution: Use the following image:

```
cm cluster define --count 3 --image CC-Ubuntu14.04
```

3. Command: cm hadoop deploy

Error: INFO: Waiting for cluster to be accessible DEBUG: Running cmd: ansible all -m ping -u ubuntu stack-0011 | UNREACHABLE! => { "changed": false, "msg": "SSH encountered an unknown error during the connection. We recommend you re-run the command using -vvvv, which will enable SSH debugging output to help diagnose the issue", "unreachable": true }

Resolution: Use the following image:

cm cluster define -count 3 -image CC-Ubuntu14.04

4. Command: cm hadoop sync

Error: no github account associated

Resolution:

git config --global user.name "username" git config --global user.email "emailid"

5. Command: apt-get install git

Error: Permission denied when running "apt-get install git" command

Resolution: Exit from the current hadoop user using *exit* command. The user will change to 'cc' that has the root privileges. Install the package through the 'cc' user log-in back as the hadoop user to use the installed package .

The "cm refresh on" command can be used most of the times to resolve multiple errors.

12. CONCLUSION

We demonstrated the implementation of a big data based predictive analytics solution using Hadoop, Spark and Spark's MLlib machine learning algorithms. We predicted the readmission likelihood with an accuracy 90% by analyzing the data stored on HDFS. While Hadoop provided us a distributed file system for storing large amounts of data, Spark provided us big data processing capabilities and several libraries to accomplish several common processing and analytics tasks. We used MLlib library's algorithms to achieve similar results as obtained with scikit-learn library. These technologies can be used to build similar solutions for real world scenarios requiring processing and performing analytics over big data. Similar applications can be built leveraging Spark's ability to process streams of data and support for machine learning algorithms.

13. ACKNOWLEDGMENTS

This project was a part of the Big Data Software and Projects (INFO-I524) course. We would like to thank Professor Gregor von Laszewski and the associate instructors for their help and support during the course.

REFERENCES

- [1] U. M. L. Repository, "Diabetes 130-US hospitals for years 1999-2008 Data Set," Web Page, accessed: 2017-03-12. [Online]. Available: <https://archive.ics.uci.edu/ml/datasets/Diabetes+130-US+hospitals+for+years+1999-2008#>
- [2] U. D. of Health and H. S. (HHS), "ICD9Data.com," Web Page, accessed: 2017-04-21. [Online]. Available: <http://www.icd9data.com/>
- [3] B. Strack, J. P. DeShazo, C. Gennings *et al.*, "Impact of hba1c measurement on hospital readmission rates: Analysis of 70,000 clinical database patient records," *BioMed Research International*, vol. 2014, no. 781670, p. 11, Apr. 2014, published as an Article. [Online]. Available: <https://www.hindawi.com/journals/bmri/2014/781670/>
- [4] Databricks, "Apache Spark," Web Page, accessed: 2017-04-21. [Online]. Available: <https://databricks.com/spark/about>
- [5] Apache, "Apache Spark MLlib," Web Page, accessed: 2017-04-21. [Online]. Available: <http://spark.apache.org/mllib/>
- [6] Apache, "Class RDD<T>," Web Page, accessed: 2017-04-21. [Online]. Available: <https://spark.apache.org/docs/1.6.2/api/java/org/apache/spark/rdd/RDD.html>
- [7] Apache, "Class RDD<T>," Web Page, accessed: 2017-04-21. [Online]. Available: <https://spark.apache.org/docs/2.0.2/api/java/org/apache/spark/SparkContext.html>
- [8] Apache, "Welcome to Apache™ Hadoop®!" Web Page, accessed: 2017-03-12. [Online]. Available: <http://hadoop.apache.org/>
- [9] Apache, "Apache Spark: Lightning-fast cluster computing," Web Page, accessed: 2017-03-12. [Online]. Available: <http://spark.apache.org/>
- [10] P. S. Foundation, "python," Web Page, accessed: 2017-03-12. [Online]. Available: <https://www.python.org/>
- [11] D. Cournapeau, "scikit-learn," Web Page, accessed: 2017-04-21. [Online]. Available: <http://scikit-learn.org/stable/>
- [12] T. Preston-Werner, "GitHub," Web Page, accessed: 2017-04-21. [Online]. Available: <https://github.com/>
- [13] R. Hat, "ANSIBLE," Web Page, accessed: 2017-03-12. [Online]. Available: <https://www.ansible.com/>
- [14] C. Cloud, "Chameleon," Web Page, accessed: 2017-04-21. [Online]. Available: <https://www.chameleoncloud.org/>
- [15] I. University, "Jetstream," Web Page, accessed: 2017-04-21. [Online]. Available: <https://jetstream-cloud.org/>
- [16] Oracle, "VirtualBox," Web Page, accessed: 2017-04-21. [Online]. Available: <https://www.virtualbox.org/wiki/VirtualBox>
- [17] LATEX, "The LATEX Project," Web Page, accessed: 2017-03-12. [Online]. Available: <https://www.latex-project.org/>
- [18] Apache, "Machine Learning Library (MLlib) Guide," Web Page, accessed: 2017-04-21. [Online]. Available: <http://spark.apache.org/docs/latest/ml-guide.html>
- [19] W. Inc., "GitHub," Web Page, accessed: 2017-04-21. [Online]. Available: <https://en.wikipedia.org/wiki/GitHub>
- [20] scikit-learn developers, "sklearn.feature_selection.VarianceThreshold," Web Page, accessed: 2017-04-21. [Online]. Available: http://scikit-learn.org/stable/modules/generated/sklearn.feature_selection.VarianceThreshold.html